

Enhanced Design Document for Distributor Data Extraction Ingestion

Nathan Lunceford

2025-02-25

Table of contents

1	Overview	2
2	System Architecture	2
2.1	Entry Point and Application Hosting	2
2.2	Extraction Ingestion Orchestration	3
2.3	External Dependencies	3
3	Design Patterns and Advanced Resilience Strategies	4
3.1	Dependency Injection (DI)	4
3.2	Registry Pattern	4
3.3	Abstract Factory Pattern	5
3.4	Builder Pattern	5
3.5	Strategy Pattern	6
3.6	Decorator Pattern	6
3.7	Circuit Breaker Pattern	7
3.8	Bulkhead Pattern	7
3.9	Exponential Backoff with Jitter	7
4	Detailed Component Descriptions	8
4.1	Program Class	8
4.2	ERPService Class	8
4.3	Providers	9
4.4	Data Uploader – S3DataUploader	9
4.5	Data Quality and Validation	10
4.6	HashiCorp Vault Integration	10
4.7	FerretDB Integration	11
4.8	Feature Flag Management	12

4.9	Metrics and Observability	12
4.10	Data Lineage and Governance	13
4.11	Configuration Models	13
4.12	Development and Test Support	14
5	Pseudo-code for Nomad Integration	14
6	Advanced Error Handling and Recovery	23
7	Security and Compliance	23
8	Development and Operations Support	24
9	Conclusion	25

1 Overview

This design document covers the extraction ingestion component of our ETL pipeline. In our architecture, distributors integrate with our system via two primary methods:

- **Direct Extraction Ingestion:** Where our system automatically provisions the correct resources based solely on the three supplied parameters: `--erp-type`, `--client-id`, and `--data-type`.
- **SFTP Ingestion:** Where distributors SFTP their data to our environment (this method is handled separately).

The design detailed here pertains only to the **direct extraction ingestion** method. This process extracts distributor data and deposits it into a pre-dropzone S3 bucket. At this stage, a minor transformation is applied: column names are standardized according to our global data dictionary and the output files are converted to Parquet format. More complex transformations occur later in the pipeline.

2 System Architecture

2.1 Entry Point and Application Hosting

- **Program Class:**
 - **Main Method:**
 - * Parses command-line options (e.g., `--erp-type`, `--client-id`, `--data-type`) using `System.CommandLine`.
 - * Sets up the dependency injection container and configures services.

- * Resolves the primary extraction ingestion service (**ERPService**) from the DI container and triggers the extraction process.
- **CreateHostBuilder Method:**
 - * Configures the host with HashiCorp Vault, FerretDB services, core providers, registries, and builder components.
 - * Configures **IHttpClientFactory** with resilience policies for API-based integrations.
 - * Sets up feature flags, metrics collection, and health checks.

2.2 Extraction Ingestion Orchestration

- **ERPService Class:**
 - Acts as the orchestrator for the extraction ingestion process.
 - Retrieves credentials from HashiCorp Vault and configuration from FerretDB, respectively.
 - Dynamically resolves components for data extraction based on ERP type.
 - Uses **IHttpClientFactory** for managing HTTP connections in API mode.
 - Supports two extraction modes:
 - * **API Mode:** For ERPs that expose APIs.
 - * **Database Mode:** For ERPs that require direct database access.
 - Supports both batch processing and incremental extraction with change data capture (CDC).
 - Applies a minimal transform:
 - * Standardizes column names according to our global data dictionary.
 - * Applies data quality validations against predefined rules.
 - * Converts all files to Parquet format with optional compression.
 - Deposits the extracted (and minimally transformed) data into a pre-dropzone S3 bucket.
 - Tracks data lineage for audit and compliance purposes.
 - (Note: Subsequent, more complex transformations are performed later in the overall ETL pipeline.)

2.3 External Dependencies

- **Infrastructure Services:**
 - **IVaultService:** Retrieves ERP credentials from HashiCorp Vault with least privilege access.
 - **IFerretDBService:** Fetches ERP configuration data from FerretDB.
 - **IAmazonS3:** Uploads the extracted data into the pre-dropzone S3 bucket.

- **IFeatureFlagService:** Manages feature flags for gradual rollout of functionality.
- **ICatalogService:** Integrates with data catalog for metadata management.
- **.NET Libraries:**
 - **System.CommandLine:** For command-line parsing.
 - **Microsoft.Extensions.Hosting & DI:** For hosting and dependency injection.
 - **System.Text.Json:** For JSON serialization/deserialization.
 - **IHttpClientFactory:** For managing HttpClient instances in API extraction mode.
 - **VaultSharp:** For interacting with HashiCorp Vault.
 - **MongoDB.Driver:** For FerretDB interactions (FerretDB uses MongoDB wire protocol).
 - **Polly:** For resilience patterns including circuit breakers and bulkheads.
 - **OpenTelemetry:** For distributed tracing and metrics.
 - **FluentValidation:** For data contract validation.

3 Design Patterns and Advanced Resilience Strategies

3.1 Dependency Injection (DI)

- **Usage:**
 - Decouples service construction from business logic.
 - Registers Vault client, FerretDB connection, providers, registries, builder components, and the main extraction service in the DI container.
 - Configures **IHttpClientFactory** and related services.
- **Benefits:**
 - Enhances testability and maintainability.
 - Promotes separation of concerns.

3.2 Registry Pattern

- **Components:**
 - **ERPRegistry**, **ExtractorRegistry**, **TransformationRegistry**, and **UploaderRegistry**.
- **Usage:**
 - Centralizes lookup for ERP-specific factories and strategies.
 - Dynamically resolves connectors, extractors, transformers, and uploaders based on ERP type or data type.
- **Benefits:**

- Simplifies addition of new ERP integrations.
- Reduces direct dependencies between the extraction process and concrete implementations.

3.3 Abstract Factory Pattern

- **Component:**
 - The `IERPFactory` interface (managed via `ERPRegistry`).
 - `IHttpClientFactory` for HTTP client management.
- **Usage:**
 - Encapsulates creation of ERP-specific components (e.g., connectors and jobs).
- **Benefits:**
 - Supports multiple ERP systems with varying implementations without altering extraction logic.

3.4 Builder Pattern

- **Components:**
 - **APIRequestBuilder & AuthenticationBuilder:**
 - * Provide fluent interfaces to construct complex API request objects.
 - **DatabaseQueryBuilder:**
 - * Dynamically constructs SQL queries for ERPs requiring direct database access.
 - **ExtractConfigBuilder:**
 - * Builds extraction configurations supporting both full and incremental extracts.
- **Usage:**
 - The `DatabaseQueryBuilder` collects parameters and produces a `DatabaseQuery` object with a `GenerateSql` method.
 - The `APIRequestBuilder` works with `IHttpClientFactory` to construct properly configured HTTP requests.
 - The `ExtractConfigBuilder` creates configurations for different extraction modes.
- **Benefits:**
 - Enhances readability and modularity.
 - Supports multiple extraction modes seamlessly.

3.5 Strategy Pattern

- **Components:**
 - Interfaces such as `IExtractor`, `ITransformer`, and `IValidator`.
- **Usage:**
 - Encapsulate different implementations for data extraction, transformation, and validation.
 - Registries select the appropriate strategy at runtime.
- **Benefits:**
 - Provides flexibility to extend or change algorithms without impacting the overall system.

3.6 Decorator Pattern

- **Components:**
 - **VaultCredentialProviderDecorator:**
 - * Adds caching behavior to Vault credential retrieval.
 - **MetricsDecorator:**
 - * Wraps core services to collect performance metrics.
 - **EncryptionDecorator:**
 - * Adds field-level encryption for sensitive data.
 - **DataQualityDecorator:**
 - * Adds data quality checks to transformations.
- **Usage:**
 - Enhances existing components with additional functionality without modifying their core logic.
- **Benefits:**
 - Improves performance through caching.
 - Provides observability and enhances security.
 - Ensures data quality through validation.

3.7 Circuit Breaker Pattern

- **Components:**
 - `VaultCircuitBreaker`, `FerretDBCircuitBreaker`, `S3CircuitBreaker`
- **Usage:**
 - Prevents cascading failures by breaking the circuit when dependencies fail.
 - Automatically restores service when dependencies recover.
- **Benefits:**
 - Enhances system resilience during partial outages.
 - Prevents overwhelming failing services with requests.

3.8 Bulkhead Pattern

- **Components:**
 - Separate connection pools for different ERP types and operations.
 - Isolated resource pools for critical vs. non-critical operations.
- **Usage:**
 - Isolates failures to prevent system-wide degradation.
 - Allocates resources based on operation criticality.
- **Benefits:**
 - Prevents resource exhaustion.
 - Improves system stability during partial failures.

3.9 Exponential Backoff with Jitter

- **Usage:**
 - Implements intelligent retry strategies for all external service calls.
 - Adds randomization to prevent thundering herd problems.
- **Benefits:**
 - Prevents overwhelming recovering services.
 - Distributes retry attempts evenly over time.

4 Detailed Component Descriptions

4.1 Program Class

- **Main Method:**
 - Parses Nomad-supplied command-line arguments.
 - Builds the DI container via `CreateHostBuilder`.
 - Resolves and invokes `ERPService.ProcessERPData`.
 - Handles global errors for graceful failure.
- **CreateHostBuilder Method:**
 - Configures services including Vault client, FerretDB connection, providers, registries, and builder components.
 - Sets up resilience policies with circuit breakers and bulkheads.
 - Configures feature flags, metrics collection, and health checks.
 - Sets up OpenAPI documentation generation.

4.2 ERPService Class

- **Responsibilities:**
 - Orchestrates the extraction ingestion process.
 - Retrieves credentials and ERP configuration.
 - Dynamically resolves ERP-specific components using registries.
 - Chooses between API or Database extraction modes based on the ERP configuration.
 - Supports both batch and incremental extraction modes.
 - Applies a minor transformation with data quality validation.
 - Deposits the transformed data into a pre-dropzone S3 bucket.
 - Tracks data lineage for audit and compliance.
- **Key Method – ProcessERPData:**
 - **Credential & Configuration Retrieval:**
 - * Uses `ICredentialProvider` to obtain short-lived, least-privilege credentials.
 - * Uses `IConfigurationProvider` to obtain ERP settings.
 - **Dynamic Component Resolution:**
 - * Uses registries to resolve ERP-specific factories, extractors, transformers, and uploaders.
 - **Integration Modes:**
 - * **API Mode:**

- Builds an API request via `APIRequestBuilder` with mutual TLS if supported.
- Extracts data via `IExtractor.Extract` with circuit breaker protection.
- * **Database Mode:**
 - Builds a SQL query using `DatabaseQueryBuilder`.
 - Extracts data via `IExtractor.ExtractFromDatabase` with connection isolation.
- **Extraction Modes:**
 - * **Full Extract:** Retrieves all data for the specified type.
 - * **Incremental Extract:** Retrieves only changed data since the last extraction.
- **Self-Healing:**
 - * Implements automatic recovery procedures for common failure scenarios.
- **Subsequent Steps:**
 - * Applies the minor transform with data quality validation.
 - * Records data lineage metadata.
 - * Uploads the resulting data into a pre-dropzone S3 bucket via `IDataUploader`.

4.3 Providers

- **VaultCredentialProvider:**
 - Retrieves and deserializes credentials from HashiCorp Vault.
 - Implements caching with TTL-based invalidation.
 - Supports dynamic secret rotation with configurable TTL.
 - Generates least-privilege, operation-specific credentials.
- **FerretDBConfigProvider:**
 - Fetches ERP configuration from FerretDB and maps it to an `ERPConfiguration` object.
 - Uses MongoDB driver since FerretDB implements MongoDB wire protocol.
 - Supports read-preference strategies for replica sets.
 - Implements schema evolution for backward compatibility.

4.4 Data Uploader – `S3DataUploader`

- **Responsibilities:**
 - Formats and uploads the minimally transformed data into S3.
 - Applies data compression for efficient storage and transfer.
 - Converts files to Parquet format and ensures standardized column names.
 - Encrypts sensitive fields before upload.

- Records data lineage metadata.
- **Key Methods:**
 - **Upload:** Manages the upload process with checksums and integrity verification.
 - **FormatData:** Applies the transformation with data quality checks.
 - **TrackLineage:** Records data provenance information.

4.5 Data Quality and Validation

- **DataContractValidator:**
 - **Responsibilities:**
 - * Validates data against predefined schemas and rules.
 - * Reports quality issues with detailed diagnostics.
 - * Enforces data governance policies.
 - **Key Methods:**
 - * **ValidateSchema:** Ensures data adheres to expected structure.
 - * **ValidateValues:** Checks data values against business rules.
 - * **GenerateReport:** Creates detailed validation reports.
- **DataMaskingService:**
 - **Responsibilities:**
 - * Identifies and masks sensitive information (PII).
 - * Supports various masking techniques (hashing, tokenization, etc.).
 - * Maintains referential integrity across masked datasets.
 - **Key Methods:**
 - * **IdentifySensitiveFields:** Automatically detects potential PII.
 - * **ApplyMasking:** Applies appropriate masking techniques.
 - * **VerifyMasking:** Ensures masking effectiveness.

4.6 HashiCorp Vault Integration

- **VaultService:**
 - **Responsibilities:**
 - * Manages connection to HashiCorp Vault with circuit breaker protection.
 - * Retrieves secrets with proper authentication.
 - * Handles secret versioning and rotation.
 - * Generates dynamic, least-privilege credentials.
 - **Key Methods:**

- * **GetSecret:** Retrieves a secret by path.
- * **GetDynamicSecret:** Retrieves a dynamic secret with lease management.
- * **RenewToken:** Handles token renewal to maintain access.
- **Configuration:**
 - * Support for multiple authentication methods (AppRole, Token, K8s).
 - * Automatic token renewal.
 - * Secret caching with TTL-based invalidation.
 - * Mutual TLS for secure communication.
- **VaultConfiguration:**
 - Stores settings such as **Address**, **AuthMethod**, **RoleId**, **SecretId**, and authentication paths.
 - Includes retry and timeout settings with exponential backoff and jitter.

4.7 FerretDB Integration

- **FerretDBService:**
 - **Responsibilities:**
 - * Manages connection to FerretDB using MongoDB driver.
 - * Executes queries and maps results to domain models.
 - * Handles connection pooling and resilience.
 - * Supports schema evolution for backward compatibility.
 - **Key Methods:**
 - * **GetConfiguration:** Retrieves configuration by ERP type and client ID.
 - * **UpdateConfiguration:** Updates configuration documents.
 - * **ExecuteQuery:** Executes a MongoDB query and returns the results.
 - * **GetSchemaVersion:** Retrieves schema version information.
 - **Configuration:**
 - * Support for replica sets and read preferences.
 - * Connection pooling with bulkhead isolation.
 - * Automatic retry with exponential backoff and jitter.
 - * Circuit breaker protection against cascading failures.
- **FerretDBConfiguration:**
 - Stores settings such as **ConnectionString**, **Database**, **Collection**, and authentication credentials.
 - Includes timeout and retry settings.
 - Defines bulkhead configuration for connection isolation.

4.8 Feature Flag Management

- **FeatureFlagService:**
 - **Responsibilities:**
 - * Manages feature flags for gradual rollout of new features.
 - * Supports A/B testing and canary releases.
 - * Provides runtime configuration without deployment.
 - **Key Methods:**
 - * **IsFeatureEnabled:** Checks if a feature is enabled for a specific context.
 - * **GetFeatureConfiguration:** Retrieves configuration for an enabled feature.
 - * **RecordFeatureUsage:** Tracks feature usage for analytics.
 - **Configuration:**
 - * Support for user, client, and global targeting.
 - * Time-based and percentage-based rollouts.
 - * Integration with monitoring for impact assessment.

4.9 Metrics and Observability

- **MetricsService:**
 - Collects performance metrics on key operations.
 - Integrates with Prometheus for metrics collection.
 - Supports custom dimensions for detailed analysis.
 - Records histogram metrics for latency distribution.
- **HealthCheckService:**
 - Provides health status of dependencies (Vault, FerretDB, S3).
 - Implements circuit breaker pattern for degraded services.
 - Exposes health endpoints for monitoring.
 - Supports self-healing procedures for common issues.
- **TracingService:**
 - Provides distributed tracing across system components.
 - Integrates with OpenTelemetry for standardized telemetry.
 - Correlates logs, metrics, and traces for holistic observability.
 - Supports sampling strategies for high-volume production environments.

4.10 Data Lineage and Governance

- **DataLineageService:**
 - **Responsibilities:**
 - * Tracks data origin, transformations, and destinations.
 - * Records metadata about extraction processes.
 - * Supports audit requirements and compliance verification.
 - **Key Methods:**
 - * **StartLineageRecord:** Creates a new lineage tracking record.
 - * **RecordTransformation:** Logs applied transformations.
 - * **CompleteLineageRecord:** Finalizes the lineage record.
 - **Integration:**
 - * Connects with data catalog for metadata management.
 - * Provides lineage visualization through API.
 - * Supports data governance and compliance reporting.

4.11 Configuration Models

- **ERPConfiguration:**
 - Stores settings such as **BaseUrl**, **CompanyId**, **WarehouseId**, **RequiredHeaders**, and timeout/retry settings.
 - **Fields for Database Access:**
 - * **AccessType:** Enum (API or Database).
 - * **ConnectionString:** For direct database connections.
 - * **Schema:** Database schema.
 - * **BatchSize:** Number of records to fetch per batch.
 - **Fields for Incremental Extract:**
 - * **SupportsCDC:** Whether the ERP supports Change Data Capture.
 - * **CDCConfiguration:** Settings for CDC-based extraction.
 - * **WatermarkColumn:** Column used for incremental extraction.
- **ERPCredentials:**
 - Holds secure API keys and client secrets.
 - Includes operation-specific scoped credentials.
 - Supports dynamic credential generation.
- **UploadConfiguration:**
 - Used by the data uploader to configure the S3 upload.
 - Includes compression settings and encryption options.

- Defines metadata for data catalog integration.
- **ResilienceConfiguration:**
 - Defines circuit breaker thresholds and recovery periods.
 - Configures retry policies with exponential backoff and jitter.
 - Specifies bulkhead isolation settings for resource pools.

4.12 Development and Test Support

- **LocalDevelopmentEnvironment:**
 - Provides containerized dependencies (Vault, FerretDB, S3-compatible storage).
 - Supports mock ERP implementations for testing.
 - Includes sample datasets for development.
- **IntegrationTestHarness:**
 - Facilitates automated testing against mock ERPs.
 - Supports scenario-based testing of the extraction process.
 - Includes performance benchmarking capabilities.

5 Pseudo-code for Nomad Integration

```

/// <summary>
/// Entry point for the ERP data extraction and ingestion process.
/// Processes command-line arguments from Nomad and orchestrates the ETL workflow.
/// </summary>
Main:
    // Parse Nomad-supplied command-line arguments:
    //   --erp-type, --client-id, --data-type
    options = parseArguments(["--erp-type", "--client-id", "--data-type"])

    // Build the host container with dependency injection configured
    host = createHostBuilder().build()

    // Register health checks for dependencies
    host.registerHealthChecks()
        .addCheck<VaultHealthCheck>("vault")
        .addCheck<FerretDBHealthCheck>("ferretdb")
        .addCheck<S3HealthCheck>("s3")

```

```

// Setup OpenAPI documentation
host.configureOpenApi(options => {
    options.Title = "Distributor Data Extraction API";
    options.Version = "v1";
    options.Description = "API for extracting distributor data from various ERP systems"
})

// Retrieve the ERPService from the DI container
erpService = host.getService(ERPService)

// Trigger the extraction ingestion process with the supplied parameters
erpService.ProcessERPData(options.erpType, options.clientId, options.dataType)

/// <summary>
/// Configures and builds the host with all necessary services.
/// </summary>
CreateHostBuilder:
return HostBuilder()
    .ConfigureServices(services =>
    {
        // Configure Vault client with circuit breaker
        services.AddSingleton<IVaultClient>(provider =>
        {
            var vaultOptions = new VaultClientSettings(
                "https://vault.example.com:8200",
                new AppRoleAuthMethodInfo(roleId, secretId)
            );
            vaultOptions.RetrySettings = new RetrySettings {
                Enabled = true,
                MaxAttempts = 5,
                BackoffType = BackoffType.ExponentialWithJitter
            };
            return new VaultClient(vaultOptions);
        });

        // Configure FerretDB connection with bulkhead isolation
        services.AddSingleton<IMongoClient>(provider =>
        {
            var settings = MongoClientSettings.FromConnectionString(
                "mongodb://ferretdb.example.com:27017"
            );
            settings.RetryWrites = true;
        });
    });

```

```

        settings.RetryReads = true;
        settings.ServerSelectionTimeout = TimeSpan.FromSeconds(5);
        settings.MaxConnectionPoolSize = 100;
        return new MongoClient(settings);
    });

    // Configure feature flag service
    services.AddSingleton<IFeatureFlagService, FeatureFlagService>();

    // Register providers with caching decorators
    services.AddSingleton<ICredentialProvider, VaultCredentialProvider>();
    services.Decorate<ICredentialProvider, CachedCredentialProviderDecorator>();

    services.AddSingleton<IConfigurationProvider, FerretDBConfigProvider>();

    // Register data quality and validation services
    services.AddSingleton<IDataContractValidator, DataContractValidator>();
    services.AddSingleton<IDataMaskingService, DataMaskingService>();

    // Register data lineage service
    services.AddSingleton<IDataLineageService, DataLineageService>();

    // Register registries
    services.AddSingleton<IERPRegistry, ERPRegistry>();
    services.AddSingleton<IExtractorRegistry, ExtractorRegistry>();
    services.AddSingleton<ITransformationRegistry, TransformationRegistry>();
    services.AddSingleton<IUploaderRegistry, UploaderRegistry>();

    // Register builders
    services.AddSingleton<IAPIRequestBuilder, APIRequestBuilder>();
    services.AddSingleton<IDatabaseQueryBuilder, DatabaseQueryBuilder>();
    services.AddSingleton<IAuthenticationBuilder, AuthenticationBuilder>();
    services.AddSingleton<IExtractConfigBuilder, ExtractConfigBuilder>();

    // Register S3 client for data upload
    services.AddAWSService<IAmazonS3>();

    // Register core service with decorators for cross-cutting concerns
    services.AddSingleton<ERPService>();
    services.Decorate<ERPService, MetricsERPServiceDecorator>();
    services.Decorate<ERPService, DataQualityDecorator>();
    services.Decorate<ERPService, EncryptionDecorator>();

```



```

// Configure HTTP clients with resilience policies using Polly
services.AddHttpClient("default")
    .AddTransientHttpErrorPolicy(builder =>
        builder.WaitAndRetryAsync(
            retryCount: 3,
            sleepDurationProvider: retryAttempt =>
                TimeSpan.FromSeconds(Math.Pow(2, retryAttempt)) +
                TimeSpan.FromMilliseconds(new Random().Next(0, 1000)), // Jitter
            onRetry: (outcome, timespan, retryAttempt, context) => {
                // Log retry attempt
                logger.LogWarning($"Retry {retryAttempt} for {context.PolicyKey}")
            }
        ))
    .AddCircuitBreakerPolicy(builder =>
        builder.CircuitBreakerAsync(
            handledEventsAllowedBeforeBreaking: 5,
            durationOfBreak: TimeSpan.FromSeconds(30),
            onBreak: (outcome, breakDelay) => {
                logger.LogError($"Circuit broken for {breakDelay.TotalSeconds}s!")
            },
            onReset: () => {
                logger.LogInformation("Circuit reset!");
            }
        ));

// Add OpenTelemetry tracing
services.AddOpenTelemetryTracing(builder => {
    builder
        .SetResourceBuilder(ResourceBuilder.CreateDefault().AddService("erp-extractor"))
        .AddSource("erp-extractor")
        .AddHttpClientInstrumentation()
        .AddMongoDBInstrumentation()
        .AddAspNetCoreInstrumentation()
        .AddJaegerExporter();
});

/// <summary>
/// Processes ERP data extraction and performs initial transformation.
/// </summary>
/// <param name="erpType">The type of ERP system to extract from</param>
/// <param name="clientId">The client identifier</param>

```

```

/// <param name="dataType">The type of data to extract</param>
/// <remarks>
/// This method handles both API and Database extraction modes. For API mode,
/// it constructs appropriate API requests with authentication. For Database mode,
/// it builds and executes SQL queries. In both cases, the extracted data is:
/// 1. Minimally transformed (column standardization)
/// 2. Validated against data contracts
/// 3. Converted to Parquet format with compression
/// 4. Uploaded to a pre-dropzone S3 bucket
/// </remarks>
ERPService.ProcessERPData(erpType, clientId, dataType):
    Log "Starting ETL extraction ingestion for client [clientId] using ERP [erpType]"

    /// <summary>Start metrics collection for this operation</summary>
    using (metricsTimer = MetricsService.StartTimer("erp_process_data",
                                                    { "erp_type": erpType, "client_id": clientId }))
    using (tracer = TracingService.StartTrace("ProcessERPData"))
    {
        /// <summary>Start data lineage tracking</summary>
        lineage = DataLineageService.StartLineageRecord(erpType, clientId, dataType)

        /// <summary>Check feature flags for enabled features</summary>
        bool useIncrementalExtract = FeatureFlagService.IsFeatureEnabled("IncrementalExtract", clientId)
        bool useCompression = FeatureFlagService.IsFeatureEnabled("Compression", clientId)
        bool useFieldEncryption = FeatureFlagService.IsFeatureEnabled("FieldEncryption", clientId)

        /// <summary>Retrieve least-privilege credentials from HashiCorp Vault</summary>
        credentials = CredentialProvider.GetLeastPrivilegeCredentials(erpType, clientId, dataType)

        /// <summary>Retrieve configuration from FerretDB</summary>
        erpConfig = ConfigurationProvider.GetConfiguration(erpType, clientId)

        /// <summary>
        /// Lookup common components via registries:
        /// - ERP-specific factory (for connectors and jobs)
        /// - Data extractor (for API or DB extraction)
        /// - Data transformer (to standardize columns and convert to Parquet)
        /// - Data validator (to validate data quality)
        /// - Data uploader (to upload data to the pre-dropzone S3 bucket)
        /// </summary>
        factory = ERPRegistry.GetFactory(erpType, clientId)
        extractor = ExtractorRegistry.GetExtractor(erpType)
    }

```

```

transformer = TransformationRegistry.GetStrategy(erpType, dataType)
validator = ValidatorRegistry.GetValidator(erpType, dataType)
uploader = UploaderRegistry.GetUploader("s3")

/// <summary>Build extraction configuration based on mode</summary>
extractConfig = ExtractConfigBuilder.New()
    .ForERP(erpType)
    .ForClient(clientId)
    .ForDataType(dataType)
    .UseIncrementalExtract(useIncrementalExtract && erpConfig.SupportsCDC)
    .WithLastExtractTime(useIncrementalExtract ? GetLastExtractTime(erpType, clientId) : null)
    .WithBatchSize(erpConfig.BatchSize)
    .Build()

/// <summary>
/// Handle database extraction mode
/// Builds and executes SQL queries for direct database access
/// </summary>
if erpConfig.AccessType == Database then:
    queryBuilder = DatabaseQueryBuilder()
        .ForERP(erpType)
        .WithConnectionString(erpConfig.ConnectionString)
        .WithSchema(erpConfig.Schema)
        .WithTable(dataType + "_table")
        .WithColumns("id", "created_at", "data")

    if useIncrementalExtract && erpConfig.SupportsCDC:
        queryBuilder.WithWhere(erpConfig.WatermarkColumn, ">", extractConfig.LastExtractTime)
    else:
        queryBuilder.WithWhere("is_processed", false)

    query = queryBuilder
        .WithOrderBy("created_at")
        .WithLimit(erpConfig.BatchSize)
        .WithCommandTimeout(erpConfig.TimeoutSeconds)
        .Build()

    Log "Executing database query: " + query.GenerateSql()

// Use bulkhead isolation for database connection
using (bulkhead = BulkheadPolicy.Execute(erpType + "-database", () => {
    extractedData = extractor.ExtractFromDatabase(query, extractConfig)
})

```

```

        return extractedData
    )))

    /// <summary>
    /// Handle API extraction mode
    /// Constructs and executes authenticated API requests with resilience
    /// </summary>
    else:
        authBuilder = AuthenticationBuilder()
            .WithApiKey(credentials.ApiKey)
            .WithClientId(credentials.ClientId)
            .WithClientSecret(credentials.ClientSecret)

        if erpConfig.SupportsMutualTLS:
            authBuilder.WithClientCertificate(credentials.ClientCertificate)

        auth = authBuilder.Build()

        requestBuilder = APIRequestBuilder()
            .ForERP(erpType)
            .WithEndpoint(erpConfig.BaseUrl + "/api/v2/sales")
            .WithMethod(GET)
            .WithAuthentication(auth)
            .WithHeaders(erpConfig.RequiredHeaders)

        if useIncrementalExtract && erpConfig.SupportsCDC:
            requestBuilder.WithQueryParameters({
                "companyId": erpConfig.CompanyId,
                "warehouse": erpConfig.WarehouseId,
                "pageSize": erpConfig.PageSize.toString(),
                "changedSince": extractConfig.LastExtractTime.toISOString()
            })
        else:
            requestBuilder.WithQueryParameters({
                "companyId": erpConfig.CompanyId,
                "warehouse": erpConfig.WarehouseId,
                "pageSize": erpConfig.PageSize.toString()
            })

        request = requestBuilder
            .WithRetryPolicy(erpConfig.MaxRetries)
            .WithTimeout(erpConfig.TimeoutSeconds)

```

```

        .Build()

        Log "Executing API request to " + erpConfig.BaseUrl + "/api/v2/sales"

        // Execute with circuit breaker protection
        extractedData = CircuitBreakerPolicy
            .ForService("erp-api-" + erpType)
            .Execute(() => extractor.Extract(request, extractConfig))

        /// <summary>Validate data against contract</summary>
        Log "Validating extracted data against contract"
        validationResult = validator.Validate(extractedData)

        if !validationResult.IsValid:
            // Handle data quality issues based on severity
            if validationResult.HasCriticalIssues():
                throw new DataContractException(
                    "Critical data quality issues detected: " +
                    validationResult.GetCriticalIssuesSummary()
                )
            else:
                // Log warnings but continue processing
                Log "Warning: Data quality issues detected: " + validationResult.GetIssuesSummary()

        /// <summary>Transform the extracted data with data lineage tracking</summary>
        Log "Starting data transformation"
        transformedData = transformer.Transform(extractedData)
        lineage.RecordTransformation("ColumnStandardization", "Standardized columns according to schema")

        /// <summary>Apply field encryption for sensitive data if enabled</summary>
        if useFieldEncryption:
            transformedData = EncryptionService.EncryptSensitiveFields(
                transformedData,
                GetSensitiveFieldsConfig(erpType, dataType)
            )
            lineage.RecordTransformation("FieldEncryption", "Encrypted sensitive fields")

        /// <summary>Apply masking for non-production environments</summary>
        if environmentType != Production:
            transformedData = DataMaskingService.ApplyMasking(
                transformedData,
                GetDataMaskingConfig(erpType, dataType)
            )

```

```

    )
    lineage.RecordTransformation("DataMasking", "Applied data masking for non-product

    /// <summary>
    /// Configure and execute the S3 upload operation
    /// Data is stored in a pre-dropzone bucket with standardized path structure
    /// </summary>
    currentTimestamp = getCurrentTimestamp()

    uploadConfig = new UploadConfiguration(
        Bucket: "erp-data-" + clientId,
        Key: erpType + "/" + dataType + "/" + currentTimestamp + "/data.parquet",
        Format: Parquet,
        Compress: useCompression,
        CompressionType: useCompression ? "SNAPPY" : null,
        Metadata: {
            "erp_type": erpType,
            "client_id": clientId,
            "data_type": dataType,
            "extract_timestamp": currentTimestamp,
            "lineage_id": lineage.Id,
            "extract_mode": useIncrementalExtract ? "incremental" : "full",
            "record_count": transformedData.Count,
            "schema_version": "1.2"
        }
    )

    Log "Starting data upload to S3 pre-dropzone"
    uploadResult = uploader.Upload(transformedData, uploadConfig)

    /// <summary>Update data catalog with metadata</summary>
    CatalogService.UpdateDatasetMetadata(
        datasetId: erpType + "-" + clientId + "-" + dataType,
        metadata: {
            "lastUpdated": currentTimestamp,
            "recordCount": transformedData.Count,
            "fileLocation": uploadResult.Location,
            "fileSize": uploadResult.Size,
            "schemaVersion": "1.2",
            "lineageId": lineage.Id
        }
    )

```

```

    /// <summary>Complete lineage record</summary>
    lineage.SetDestination(uploadResult.Location)
    lineage.CompleteLineageRecord()

    /// <summary>Update last extract time for incremental extracts</summary>
    if useIncrementalExtract:
        StoreLastExtractTime(erpType, clientId, dataType, currentTimestamp)

    Log "Extraction ingestion process completed successfully"
}

```

6 Advanced Error Handling and Recovery

- **Error Classification:**
 - Categorizes errors as transient or persistent.
 - Applies different recovery strategies based on error type.
 - Records error patterns for proactive monitoring.
- **Self-Healing Procedures:**
 - **Common Issue Resolution:**
 - * Automatic token renewal for expired credentials.
 - * Connection pool refresh for stale connections.
 - * Temporary file cleanup for storage-related issues.
 - **Degraded Mode Operation:**
 - * Falls back to full extract if incremental extract fails.
 - * Disables optional features during high load.
 - * Implements progressive backoff for system recovery.
- **Comprehensive Logging:**
 - Structured logging with correlation IDs across components.
 - Context-enriched log entries for easier debugging.
 - Log level adjustment based on operation criticality.
 - Integration with log aggregation systems.

7 Security and Compliance

- **Least Privilege Access:**

- Dynamic generation of operation-specific credentials.
- Short-lived tokens with minimal required permissions.
- Credential scoping based on operation context.
- Role rotation and separation of duties.
- **Data Protection:**
 - Field-level encryption for sensitive data.
 - Data masking for non-production environments.
 - Secure credential management with HashiCorp Vault.
 - Mutual TLS for secure service communication.
- **Audit and Compliance:**
 - Comprehensive data lineage tracking.
 - Access and operation audit logging.
 - Compliance validation against regulatory requirements.
 - Regular security posture assessment.

8 Development and Operations Support

- **Local Development Environment:**
 - Docker Compose setup with all dependencies.
 - Mock ERPs for integration testing.
 - Configuration templates for different scenarios.
 - Development tools for data visualization and debugging.
- **Deployment and Infrastructure:**
 - GitOps-based deployment pipelines.
 - Infrastructure as Code for all components.
 - Canary deployment support for risk mitigation.
 - Blue-green deployment capability for zero downtime.
- **Documentation and Knowledge Sharing:**
 - OpenAPI documentation for all service interfaces.
 - Architecture decision records (ADRs) for design choices.
 - Runbooks for common operational procedures.
 - Automated documentation generation from code.

9 Conclusion

The extraction ingestion component of the ETL pipeline has been redesigned with a focus on resilience, security, performance, and maintainability. The system now supports both batch and incremental extraction modes, implements comprehensive data quality validation, and provides robust error handling with self-healing capabilities.

The transition from AWS Secrets Manager to HashiCorp Vault and from DynamoDB to FerretDB enhances security, scalability, and operational flexibility while maintaining compatibility with existing systems. The addition of circuit breakers, bulkheads, and exponential backoff strategies significantly improves system resilience during partial failures.

Enhanced data governance capabilities, including data lineage tracking, field-level encryption, and data masking, ensure compliance with regulatory requirements and protect sensitive information. Feature flags enable gradual rollout of new functionality and provide flexibility for operational adjustments without code changes.

The architecture and design patterns employed enable a flexible, maintainable, and scalable solution that can adapt to changing requirements and handle the varied needs of different ERP integrations at scale.